



Doing Power BI The Right Way

for Database Developers

Paul Turley
Principal Consultant,
3Cloud Solutions



Level: Intermediate

What's This All About?

From the perspective of a DBA and database professional, this session will provide prescriptive guidance to plan and build an enterprise BI solution correctly using best practice design. The mindsets of self-service BI and database development are often different, so using common language and a best-of-breed approach is a crucial step for designing futureproof BI solutions, and “doing Power BI the right way”.

Based on 25 years of industry experience and hundreds of durable Power BI solutions, you will see how to create and maintain analytic reporting projects supporting data governance, self-service reporting; and agile, iterative development practices.

Paul Turley

Principal Consultant, Microsoft Data Platform MVP

~25 years in IT, data platform,
Business Intelligence & data analytics



Agenda

- Enterprise Readiness
- Futureproofing Power BI solutions
- Decision points:
 - Managing Project Files
 - Accessing Source Data
 - Object Naming
 - Model Design
 - Calculations & Measures
 - Data Security
 - Deployment & Updates

Typical Power BI Use Cases

- Ad hoc analysis
- Quick reporting
- Transformation & shaping as-you-go
- Data model evolution
- Visualize: try-and-see
- Standards? Who needs standards?
 - Naming conventions
 - Modeling & relationships
 - Measure logic
 - Documentation
 - Version control



Database Development Discipline

- Data quality is enforced in the database
- Database is single version of the truth
- Avoid data copies & silos
- All requirements are documented
- Version control
- DevOps: DEV, TEST & PROD
- Scripted & automated
- Naming conventions:
PascalCase, camelCase






Audience, Size & Scale

One size doesn't always fit all

With planning, informal projects can be adapted to meet enterprise needs

Know when to keep and update a design

Know when to toss a design and start over



Power BI Pro License

1 GB max
1 license per report user

Premium Per User

100 GB
Premium Capacity
400 GB

Premium Capacity

Hybrid Data Models

No stated size limits
No licenses needed for report viewers

How Much Data?

Power BI Tabular Engine is optimized for in-memory scans & calculations

Data volume:

- Millions of rows ?
- Billions of rows ?

Data model size:

- 1 GB
- 10 GB
- 100 GB
- 400 GB
- Terabytes & Petabytes

Licensing Power BI

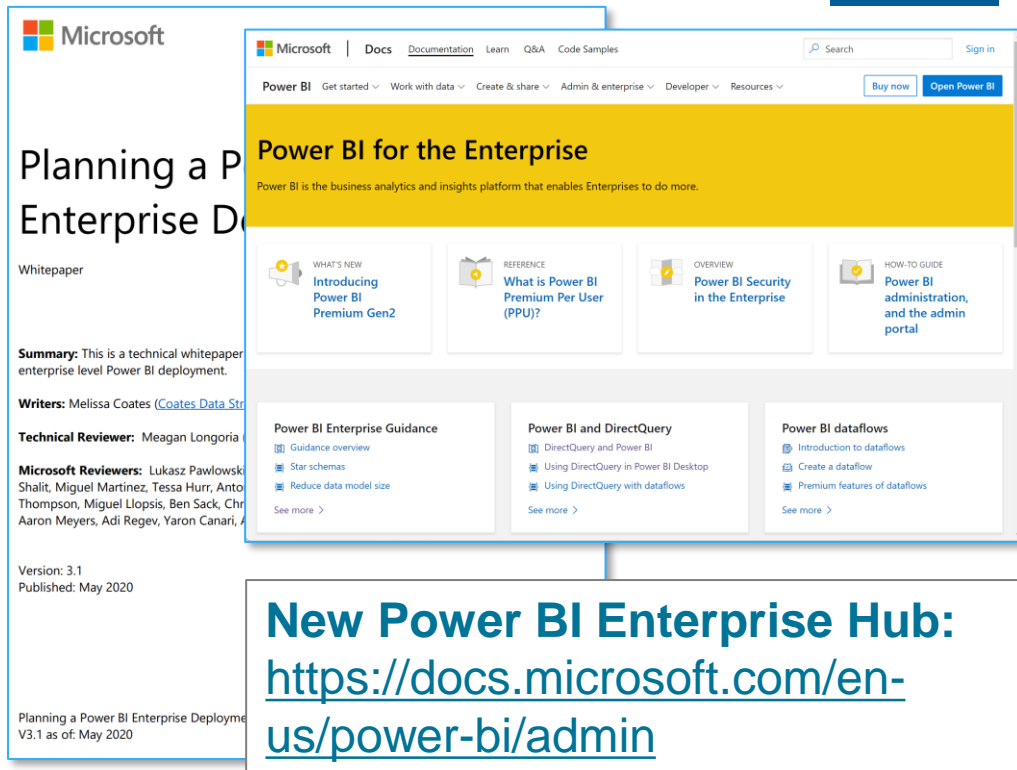
(guidelines)

Power BI Pro	\$10/user
Premium per User	\$20/user
Premium Capacity (P1)	\$5k/Month
A1 – A6	\$735 - \$3k
P2 – P5	\$10k - \$80k

Is Power BI Enterprise Ready?

- ✓ Large-scale Data
- ✓ Broad user support
- ✓ Team Development & Collaboration
- ✓ Data Governance Policies
- ✓ Data Management Policies
- ✓ Lifecycle & Deployment
- ✓ Versioning
- ✓ Embedding & Integration
- ✓ Automation
- ✓ Administration
- ✓ Security:
 - User Access & Data Protection
 - Role-based & User-based filtering

White papers for Power BI - Power BI | Microsoft Docs:
<https://docs.microsoft.com/en-us/power-bi/guidance/whitepapers>



Microsoft

Power BI for the Enterprise

Power BI is the business analytics and insights platform that enables Enterprises to do more.

WHAT'S NEW
Introducing Power BI Premium Gen2

REFERENCE
What is Power BI Premium Per User (PPU)?

OVERVIEW
Power BI Security in the Enterprise

HOW-TO GUIDE
Power BI administration, and the admin portal

Power BI Enterprise Guidance
Guidance overview
Star schemas
Reduce data model size
See more >

Power BI and DirectQuery
DirectQuery and Power BI
Using DirectQuery in Power BI Desktop
Using DirectQuery with dataflows
See more >

Power BI dataflows
Introduction to dataflows
Create a dataflow
Premium features of dataflows
See more >

Version: 3.1
Published: May 2020

Planning a Power BI Enterprise Deployment
V3.1 as of: May 2020

New Power BI Enterprise Hub:

<https://docs.microsoft.com/en-us/power-bi/admin>

Certified & Shared Datasets



- Use Dataset endorsement & certification in the service
- Certification can be managed by security group
- Access to datasets can be restricted to certified datasets
- Organization defines certification policy & provides documentation

Showing 8 items

NAME ↑	ENDORSEMENT	ACTIONS	REFRESHED	NEXT REFRESH	API A
ContosoDW Sales	Promoted		5/1/2019, 8:04:36 AM	N/A	--
ContosoDW Sales (incomplete)			3/27/2019, 10:00:36 PM	N/A	--
ContosoDW Sales (PBDT 4-19)			3/20/2019, 2:30:26 PM	N/A	--
Lightning Talk - Web API Data Sources (0...			11/7/2018, 10:27:04 AM	N/A	--
World Peace	Certified		9/25/2018, 10:45:43 PM	N/A	--
World Peace-2018-09-27			9/27/2018, 2:23:21 PM	N/A	--

Admin portal

Usage metrics
Users
Audit logs
Tenant settings
Capacity settings
Embed Codes
Organizational visuals
Dataflow settings
Workspaces

Certification
Unapplied changes
Allow users in this org to certify datasets.

☒ Enabled

Specify URL for documentation page
Enter URL

Apply to:
☐ The entire organization
☒ Specific security groups

Enter security groups

☐ Except specific security groups

Apply Cancel

Futureproofing Power BI Solutions



*Difficult to
see;
always in
motion,
the future
is.*

You will never ever, ever..

EVER EVER

have ALL requirements defined ahead of time

“Just one more thing”



Iterative Design

- Define solution scope
- Work in iterative cycles
- Evaluate design
- Keep what works
- Discard what doesn't
- Learn & improve
- Evaluate to solution scope



Database Developer



- Use source code repository
- Automated version control
- Manage everything in one place
- Automate build & deployment

BI Solution Designer



- Shared file folder
- Version control
- Desktop

Balance

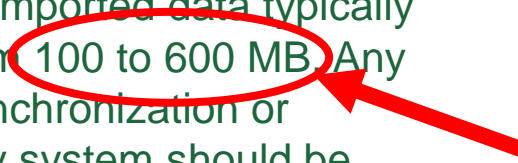
- Shared storage of some kind
- Roll-back & disaster recovery
- Manual deployment is easy
- Automated deployment is possible
- Don't merge & split PBIX files

Manage Power BI Desktop Files

- **Store files in a centrally managed network-assessable folder**
The storage folder should support automatic backup and recovery in the case of storage loss.
- **Report and dataset developers must open files from the Windows file system**
Files must either reside in or be synchronized with the Windows file system.
- Files containing imported data typically range in size from 100 to 600 MB. Any shared folder synchronization or disaster recovery system should be designed to effectively handle multiple files of this size.

Options:

- OneDrive For Business (shared by team, with folder synchronization).
- SharePoint or SharePoint Online (with folder synchronization).
- GitHub and/or VSTS with local repository & folder synchronization. If used, Git must be configured for large file storage (LFS) if PBIX files are to be stored in the repository.



We'll talk about this a bit later

Database Developer



- Write SQL
- Process queries
- Real time conn

BI Solution Designer



model
ons in-memory

Balance

- Use views to encapsulate SQL queries
- Ensure query folding works on large tables
- Use Import by default
- Use DirectQuery only when required

Transformation Options

Dataflows

Design in browser
Execute in service

In-place of data
warehouse
Reusable queries


Power Query

Design in Desktop
Execute in service

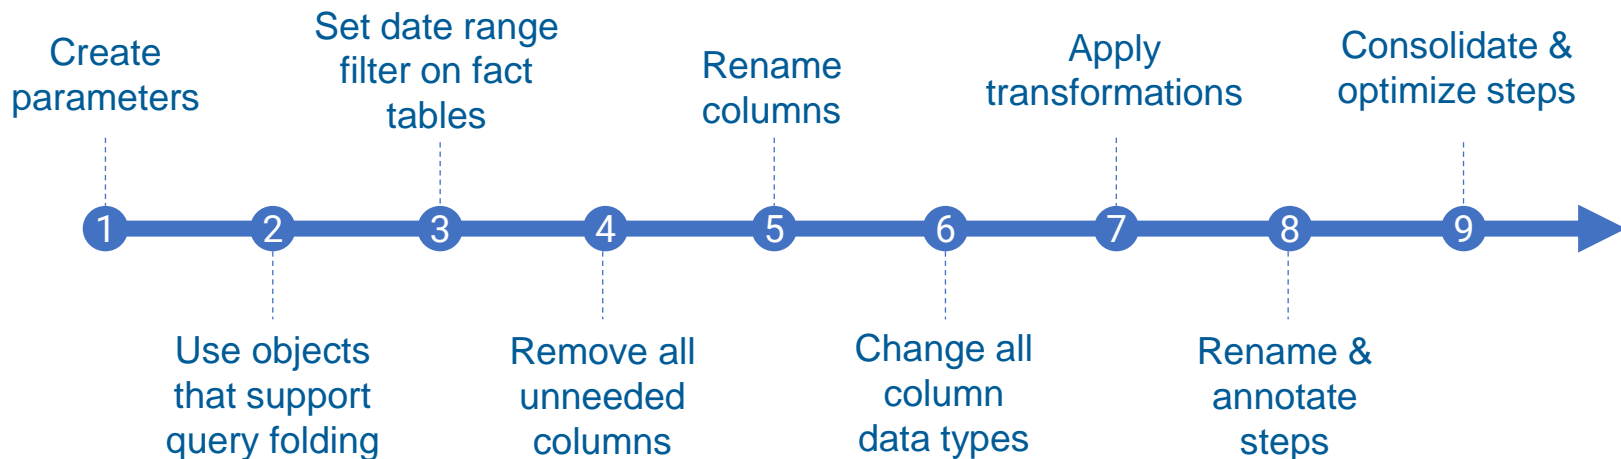
Small-moderate data
volume
Self-service data mashup

ETL/ELT & Views

Execute in database
Integrated in data
warehouse
Enterprise-scale



Preparing, shaping & transforming source data using Power Query



SQL Isn't Always The Answer



Import tables and views rather than hand-written SQL queries, so Power Query can fold queries and execute them at the source

- Prepare dimensional tables in the source database
- Create database views

Never Forget This

A close-up photograph of an elephant's head, showing its large ears, trunk, and tusks. The elephant is looking slightly to the right. The background is a clear blue sky.

**Transform Data as Far
Upstream as Possible**

Downstream transformations may be more flexible for self-service models but may not scale

Database Developer



- Table, field/column names
- Measure names
- Code-friendly names
- Cryptic object names
- Pascal case, Camel case
- Hungarian notation

BI Solution Designer



Balance

- All objects exposed to users should have friendly names
- Hide key columns & utility objects
- Hide numeric columns & create friendly named measures

Remember One Thing



Name Objects for Usability

Data model consumers are business users, not developers. Name tables, fields & measures so they are intuitive & easy to understand.

~~ACCT_TYP~~
customerPhoneNbr

Account Type
Customer Phone #

Database Developer



- Flattened result
- Transform in SQL
- Wide tables
- Pre-aggregated
- Import summary

BI Solution Designer



Power Query


Balance

- Build dimensional star schema
- Avoid wide tables
- Remove long text fields
- Remove unused fields
- Tall tables are OK if they are optimized for storage & analytics

Data modeling essentials & best practices in Power BI




Flat Model



Master/Detail





Dimensional

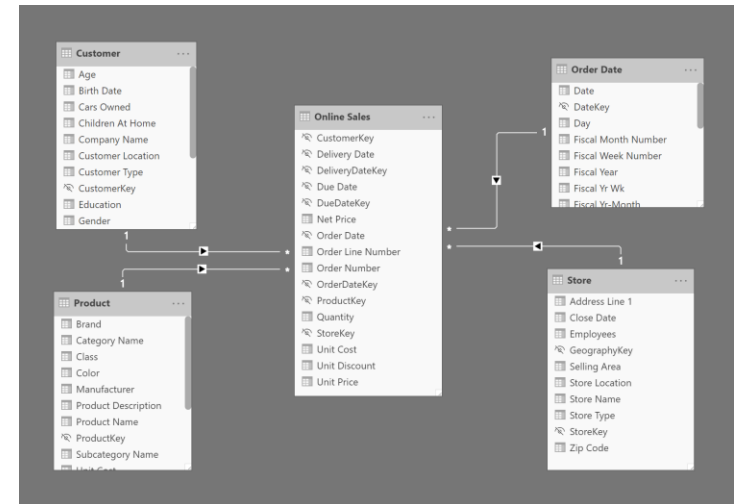


Dimensionalize It!

	EMPLOYEE [RP]	CARRIER	PRODUCT	SALES LOCATION	WAREHOUSE	PROMOTION	ORDER ID [DD]	SHIP MODE	SHIPMENT NUMBER [DD]
	who		what	where			why & how		
SALES TARGETS									
salesperson has product type target	✓		✓	✓					
CUSTOMER ORDERS									
customer orders product	*	✓	✓	✓	*		✓	*	
PRODUCT SHIPMENTS									
warehouse worker ships product	✓	✓	✓	✓	✓	✓		✓	✓
CARRIER DELIVERIES									
carrier delivers product	✓	✓	✓	✓	✓			✓	✓
CUSTOMER COMPLAINTS									
customer complains about product	✓	✓	✓	✓			✓	✓	
PRODUCT RETURNS									
customer returns product	✓		✓	✓	✓		✓	✓	



- Nearly any reporting data can be shaped into a star schema & optimized for analysis
- Ad hoc models never evolve into a dimensional model



Model Design Guidelines



- Dimensional design concepts haven't changed in 20 years & are as true as ever
- Dimensional modeling “rules” should be followed but can be relaxed for Power BI in certain cases, such as:
 - Leaving some dimensional attributes in fact tables
 - Use natural keys rather than generating surrogate keys
- The art of dimensional modeling ranges from simple to complex. Start with the basics.
- Flattened “spreadsheet” models are OK for small, informal projects but have significant limitations
- As models grow in size & complexity, data quality challenges will surface that can be solved by implementing proper governance controls

The Kimball Method: <https://www.kimballgroup.com/data-warehouse-business-intelligence-resources/kimball-techniques/dimensional-modeling-techniques>

Lawrence Corr, Model Storming Agile method: <https://modelstorming.com/hierarchy-map>

Planning for Separation – data models & reports



The Thick and Thin of Reports

Separate reports and data models can be:

- Versioned
- Deployed & managed separately
- Central “Golden” dataset
- Use Live Connect when creating reports



Doing Power BI the Right Way: 7.

Planning for separation – data models and reports | Paul Turley's SQL Server BI Blog

5 Tips for Separating Power BI Datasets and Reports — Coates Data Strategies

Hot Swap tool demo: PowerBI.tips

PowerBI.tips - Hot Swap Connections Webinar – YouTube <https://www.youtube.com/watch?v=syglUPMlgi0>

Enterprise Scale Options

In many ways, Power BI has now surpassed the capabilities of SQL Server Analysis Services. Microsoft are investing in the enterprise capabilities of the Power BI platform by enhancing Power BI Premium Capacity, adding Paginated Report and features to support massive scale specialized use cases. Consider the present and planned capabilities of the Power BI platform; before, choosing another data modeling tool such as SSAS.

Resources:

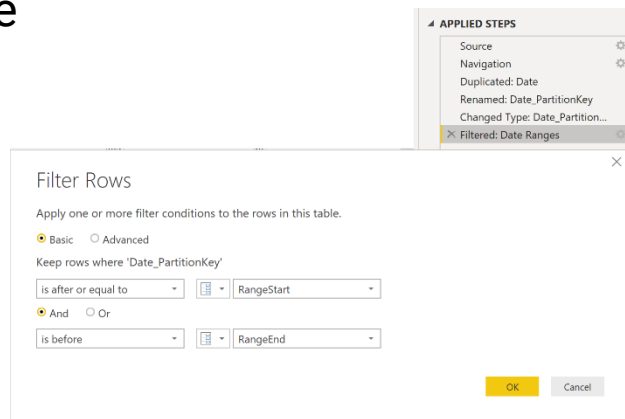

<https://sqlserverbi.blog/2018/07/27/power-bi-for-grownups>

<https://sqlserverbi.blog/2018/12/13/data-model-options-for-power-bi-solutions>

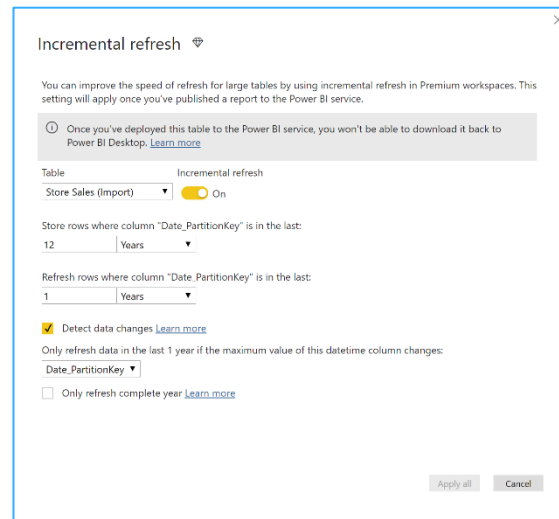
Managing Dataset Size with Parameters



- Use parameters whether implementing incremental refresh or not
- **RangeStart & RangeEnd** parameters must be date/time type
- Apply range filter on date/time column in Power Query



*Compare row count & refresh history –
1 mo & incr refresh*



***Incremental Refresh is a Premium [capacity|user] feature but parameters may be used in any data model**

Remember This



Use Parameters to Filter Results

Query parameters are used to reduce the model size in development

Implement Incremental Refresh

Scale model size & data volume

Database Developer



- Perform calculations
- Store calculated
- Process calculations in the database

BI Solution Designer



- **Balance**
- Understand & embrace row & filter context
- Store detail-level calculations
- Design measures to work with grouped fields & hierarchies

Calculations in DAX
Report context
Calculation logic in the

Database Developer



- Lock it all down
- Keep data in the
- Manage security

BI Solution Designer



data for me
in the data

Balance

- Manage privacy & security according to data governance policy
- Apply row-level security in data model
- Create & maintain user mapping table to support row-level security

Database Developer



- Rigid development
- Documented requirements
- Change requests
- Long release cycle
- Tidy/disciplined/planned

BI Solution Designer



- Right-size process for the audience & business
- Sandbox self-service models & reports in workspaces
- Apply rigid controls to certified models & reports using deployment pipelines & workspace apps

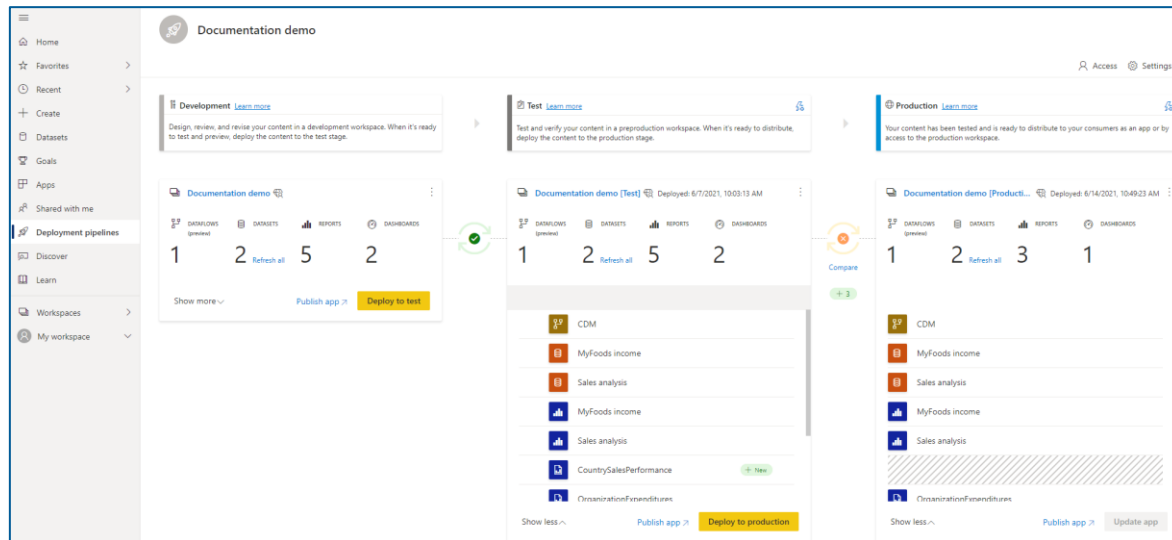
Deployment lifecycle requirements as needed
Controlled/chaotic

Balance

Deployment Pipelines

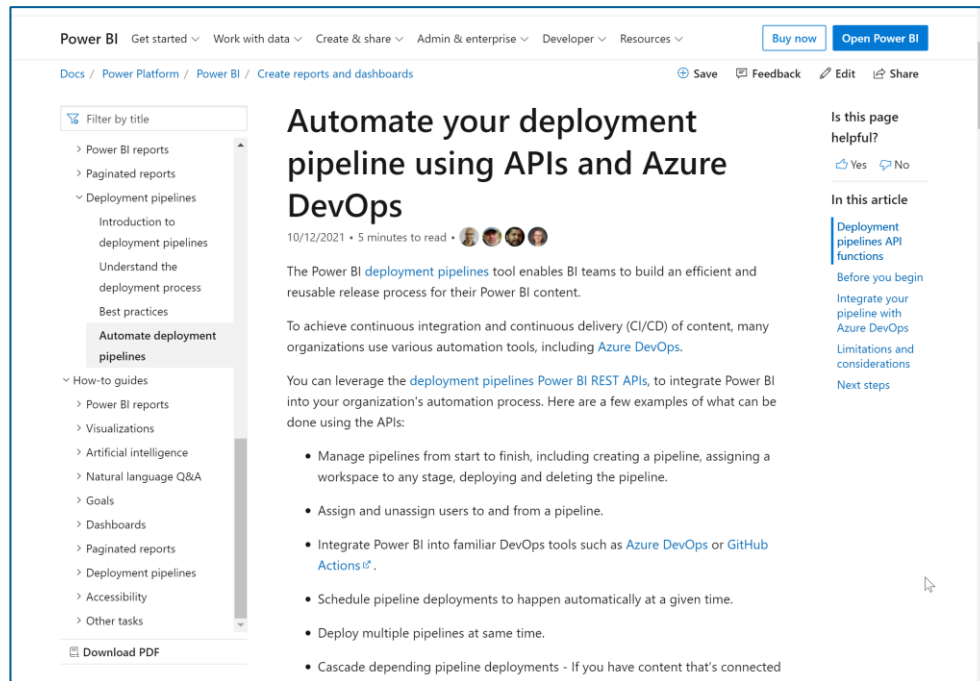


- Deployment pipelines manage the promotion of workspace app content through lifecycle stages
- Versioning with regression



DevOps Integration

- Deployment pipelines API functions:
 - Deploy all
 - Selective deploy
 - Backward deploy
 - Update App
- REST APIs may be called with PowerShell or other automation code
- Supports OAuth



<https://docs.microsoft.com/en-us/power-bi/create-reports/deployment-pipelines-automation>

Multi-developer & Lifecycle Management for Power BI



DATA MODEL

- Separate data model PBIX from Report PBIX
- Deploy and manage separately
- Migrate large-scale or enterprise data models to Model.BIM file
- Manage enterprise models in Tabular Editor or Visual Studio
- Use Deployment Pipelines for Apps

REPORTS

- Treat report PBIX files as binary files
- Use OneDrive or code repo versioning
- Don't use merging and differencing

Take Aways

- Store source files in shared storage
- Separate data model from reports divide work
- Understand and promote Query Folding
- Use SQL views rather than in-line SQL
- Use parameters to reduce working data model file size
- Consider using Incremental Refresh to manage large tables
- Use Deployment Pipelines to manage delivery

In Summary: Model Design Checklist



☐ **Model for the user experience, not for developers**

☐ **Build star schemas**

Wherever possible, reshape data into fact a dimension tables with single key, one-to-many relationships from dimensions to fact.

☐ **Enforce dimension key uniqueness**

Just because a key value “should” be unique, there is no guarantee that it will be unless enforced at the data source. Perform grouping and duplicate reduction in the data source views or Power Query queries to guarantee uniqueness. Duplicate record count checks and other mechanisms can be applied to audit source data for integrity but do not allow the data model to violate these rules.

☐ **Avoid bi-directional filters & unnecessary bridging tables**

These data modelling patterns adversely affect performance.

☐ **Consider using DAX measures** rather than complex & inefficient relationships

☐ **Create custom columns in Power Query**

Rather than DAX calculated columns wherever possible for row-level derived columns. This maintains a consistent design pattern for maintainability.

☐ **Annotate code**

Use in-line comments and annotations in all code including SQL, M and DAX; to explain calculation logic and provide author and revision information.

☐ **Remove all unused fields** – if in doubt, take it out

☐ **Hide all fields not used directly by users**

primary and foreign key columns, numeric columns used to create measures, and columns used to specify the sort order of other fields.

☐ **Use friendly field names**

Rename all visible columns (in Power Query) to short but user-friendly names with mixed case and spaces.

☐ **Set to Do Not Summarize**

Any non-hidden numeric columns that are not intended to roll-up or summarize values. Columns set to summarize are indicated with a Sigma icon.

A blue-tinted photograph of a business meeting in a modern office with large windows. Four people are gathered around a table, looking at documents. A large, semi-transparent white question mark is overlaid in the center of the image.

Questions?

- Q&A
- Discussion

Resources, Contact & Network



- **Blogs:**

- SqlServerBi.blog
- DataOnTheRoad.blog
- LinkedIn.com/in/Pturley
- @paul_turley

- **New Power BI Enterprise Hub:**

<https://docs.microsoft.com/en-us/power-bi/admin>

- **Guy In A Cube:**

<https://www.youtube.com/channel/UCFp1vaKzpfvoGai0vE5VJ0w>

- Saturday morning livestream
10:30 AM EST